

SARTHAK CHHABRA

Technical Lead · Backend & LLM Infrastructure

Gurgaon, India · +91-9557722510 · hello@sarthakchhabra.com · sarthakchhabra.com

PROFILE

Backend engineer and technical lead with 7+ years building distributed systems and production LLM infrastructure at consumer scale. Most recent work includes agent platforms, RAG pipelines and a routing engine serving 300K+ LLM requests a day. Built no-code AI tooling that non-technical teams use to ship production agents, and led teams of up to 4 engineers while owning architecture end to end.

SKILLS

Languages & Runtime: Node.js, TypeScript, JavaScript, SQL

AI / LLM Systems: LangChain, LangGraph, CrewAI, multi-agent architecture, agent orchestration, RAG pipelines, Qdrant, embeddings, LLM routing and fallbacks, prompt and tool orchestration, LLM observability (Langfuse, LangSmith, OpenTelemetry/OTel)

Data & Infra: PostgreSQL, MongoDB, Kafka, AWS (SQS, core services), Docker, microservices

Practices: System design, technical architecture, team leadership and mentoring, API design

WORK EXPERIENCE

Technical Lead · Stashfin

Feb 2026 – Jul 2026

Node.js · TypeScript · AWS · PostgreSQL · Generative AI · Docker · Microservices

- Architected and led a no-code platform for building and deploying production AI agents. Agent delivery went from days to minutes with a self-serve UI for prompts, tools, models and channels. It now runs multiple production agents, including a customer support agent handling around 30,000 messages and 44M tokens a day for 6,000+ daily users.
- Designed a tool framework that turns any REST API into an agent-callable tool. The user configures the request, the platform fires a live test call to capture the real response shape, and annotated response keys map into the LLM's tool schema. No hand-written integrations needed.
- Built a cross-channel contact system that merges a user's email, phone and Slack identity into one record with persistent memory, so an agent keeps context when the user switches channels instead of starting cold.
- Developed the RAG ingestion pipeline (upload, chunking, embedding, vector storage) that lets non-technical users give agents custom knowledge bases without engineering help.
- Shipped multi-channel delivery for Slack, Telegram and other channels, abstracting per-channel auth so a single agent configuration deploys everywhere.

Technical Lead, Gen AI · Zupee (Cashgrail Pvt Ltd)

Jul 2025 – Feb 2026

Role ended after India's 2025 online real-money gaming law shut down the company's core product.

Node.js · PostgreSQL · Qdrant · MongoDB · LangGraph · Generative AI · Microservices

- Architected the central LLM routing engine: 300,000+ requests a day at 15-20ms overhead and 99% uptime, with budget tracking, rate limiting, multi-provider fallbacks and circuit breakers.
- Developed an AI companion chatbot that reached 46% D15 retention, using LangGraph multi-agent planning, RAG memory on Qdrant, and response timing tuned to feel like human conversation.
- Built an automated promo generation system that cut production time from a week to about 3 hours by processing 50-100 microseries in parallel.
- Created a no-code bot management platform so product managers can change AI behavior, swap models and manage image catalogues themselves, in a couple of clicks.
- Led a team of 4 engineers across multiple AI products while owning the architecture, including in-house embeddings for operational control.

Senior Software Engineer · Dresma AI

Apr 2022 – Apr 2025

Node.js · TypeScript · AWS · MongoDB · Kafka · Microservices

- Owned backend architecture for heavy-computation processing pipelines on Node.js, Kafka and AWS, cutting processing time about 30% and improving stability.
- Cut API response times roughly 40% through profiling, debugging and optimization of core services.
- Built fault-tolerant services for high-load workloads: queue-backed retries on Kafka/SQS, graceful degradation, and monitoring that brought down incidents and recovery time.
- Trained and mentored junior developers on backend, frontend and AWS infrastructure, raising delivery quality and team autonomy.

Full Stack Developer · Gigforce

May 2021 – Apr 2022

Node.js · AWS SQS · MongoDB · Vue.js

- Led development of new backend modules on Node.js and MongoDB, improving stability and query performance with better indexing and service-level tuning.
- Built queue-based services on AWS SQS to absorb growing operational scale while keeping latency low at peak hours.

Software Engineer · SignCatch

Oct 2020 – May 2021

React · Angular · PHP · SQL

- Developed and maintained new and legacy modules on a React, Angular and PHP stack, taking features from requirements through to production releases.

Associate Software Engineer · Bosch

Aug 2019 – Oct 2020

Joined as a Project Trainee intern, Jan 2019 to Jul 2019, before moving to a full-time role.

Angular · C# · SQL

- Built cross-platform application features in Angular and C#, working closely with product owners to keep development aligned with product goals and cut rework.
- Triaged and resolved user-reported issues, escalating the complex ones to the right teams and reducing recurring problems through shared practices.

EDUCATION

B.E. Computer Science (8.7 CGPA) · Lovely Professional University

2015 – 2019